

Icing

Supporting Fast-Math Style Optimizations in a Verified Compiler

**Heiko Becker, Eva Darulova,
Magnus Myreen, Zachary Tatlock**



MAX PLANCK INSTITUTE
FOR SOFTWARE SYSTEMS



CHALMERS
UNIVERSITY OF TECHNOLOGY



How we develop programs



readability over performance

How we develop programs



readability over performance



compiler should make program fast

How we develop programs



readability over performance

Compiler optimizations are
a vital part of
the development process

↓
make program fast

The need for understandable optimizations

What does gcc's ffast-math actually do?

[Ask Question](#)

▲
130
▼
★
44

I understand gcc's `--ffast-math` flag can greatly increase speed for float ops, and goes outside of IEEE standards, but I can't seem to find information on what is really happening when it's on. Can anyone please explain some of the details and maybe give a clear example of how something would change if the flag was on or off?

I did try digging through S.O. for similar questions but couldn't find anything explaining the workings of ffast-math.

[performance](#)[math](#)[gcc](#)[floating-point](#)[fast-math](#)

The need for understandable optimizations

What does gcc's ffast-math actually do?

[Ask Question](#)

▲
130
▼
★
44

I understand gcc's `--ffast-math` flag can greatly increase speed for float ops, and goes outside of IEEE standards, but I can't seem to find information on what is really happening when it's on. Can anyone please explain some of the details and maybe give a clear example of how something would change if the flag was on or off?

I did try digging through S.O. for similar questions but couldn't find anything explaining the workings of ffast-math.

[performance](#)[math](#)[gcc](#)[floating-point](#)[fast-math](#)

asked 7 years, 10 months ago

viewed 41,233 times

active 10 months ago

The need for understandable optimizations

What does gcc's ffast-math actually do?

[Ask Question](#)

▲
130
▼

I understand gcc's `--ffast-math` flag can greatly increase speed for float ops, and goes outside of IEEE standards, but I can't seem to find information on what is really happening when it's on. Can anyone please explain some of the details and maybe give a clear example of how something would change if the flag was on or off?

★
44

I did try digging through S.O. for similar questions but couldn't find anything explaining the workings of ffast-math.

[performance](#)[math](#)[gcc](#)[floating-point](#)[fast-math](#)

asked 7 years, 10 months ago

viewed 41,233 times

active 10 months ago

The long road of compiler verification

1967: *McCarthy, Painter;*

Correctness of a Compiler for Arithmetic Expressions (Integer's only)

2009: *Leroy;*

Formal Verification of a Realistic Compiler

2019: *Lööw et al;*

Verified Compilation on a Verified Processor

The long road of compiler verification

1967: *McCarthy, Painter;*

Correctness of a Compiler for Arithmetic Expressions (Integer's only)

2009: *Leroy;*

Formal Verification of a Realistic Compiler

← CompCert C compiler

2019: *Lööw et al;*

Verified Compilation on a Verified Processor

The long road of compiler verification

1967: *McCarthy, Painter;*

Correctness of a Compiler for Arithmetic Expressions (Integer's only)

2009: *Leroy;*

Formal Verification of a Realistic Compiler

← CompCert C compiler

2019: *Lööw et al;*

Verified Compilation on a Verified Processor

← CakeML & Silver

The long road of compiler verification

1967: *McCarthy, Painter;*

Correctness of a Compiler for Arithmetic Expressions (Integer's only)

2009: *Leroy;*

Formal Verification of a R

Where are floating-points?

2019: *Lööw et al;*

Verified Compilation on a Verified Processor


The long road of compiler verification

1967: *McCarthy, Painter;*

Correctness of a Compiler for Arithmetic Expressions (Integer's only)

2009: *Leroy;*

Formal Verification of a Realistic Compiler



2015: *Boldo et al.;*

Verified Compilation of Floating-Point Programs

2019: *Lööw et al.;*

Verified Compilation on a Verified Processor

The state-of-the-art for fast-math

Unverified Compilers (gcc, clang,)

Verified Compilers (CakeML, ...)

The state-of-the-art for fast-math

Unverified Compilers (gcc, clang,)

Verified Compilers (CakeML, ...)

- apply aggressive optimizations

The state-of-the-art for fast-math

Unverified Compilers (gcc, clang,)

- apply aggressive optimizations
- do not preserve IEEE754 semantics

Verified Compilers (CakeML, ...)

The state-of-the-art for fast-math

Unverified Compilers (gcc, clang,)

Verified Compilers (CakeML, ...)

- apply aggressive optimizations
- do not preserve IEEE754 semantics
- give no guarantees on the result

The state-of-the-art for fast-math

Unverified Compilers (gcc, clang,)

- apply aggressive optimizations
- do not preserve IEEE754 semantics
- give no guarantees on the result

Verified Compilers (CakeML, ...)

- apply no floating-point optimizations

The state-of-the-art for fast-math

Unverified Compilers (gcc, clang,)

- apply aggressive optimizations
- do not preserve IEEE754 semantics
- give no guarantees on the result

Verified Compilers (CakeML, ...)

- apply no floating-point optimizations
- fully preserve IEEE754 semantics

The state-of-the-art for fast-math

Unverified Compilers (gcc, clang,)

- apply aggressive optimizations
- do not preserve IEEE754 semantics
- give no guarantees on the result

Verified Compilers (CakeML, ...)

- apply no floating-point optimizations
- fully preserve IEEE754 semantics
- guarantee preserving literal meaning

The state-of-the-art for fast-math

Unverified Compilers (gcc, clang,)

Verified Compilers (CakeML, ...)

- apply aggressive optimizations
- do not preserve IEEE semantics
- give no guarantees about literal meaning


We need a semantics to
handle fast-math optimizations
in verified compilers

Contributions

Icing, a **nondeterministic** semantics for floating-points:

- Support for subset of **gcc's fast-math optimizations**
- Optimization with **fine-grained control**
- Implementation of **three optimizers**
- Verification in HOL4
- Connection to CakeML

Optimizations in Icing



Example Optimizations:

$$a + b \longrightarrow b + a$$


$$a \times b \longrightarrow b \times a$$

$$a + (b + c) \longrightarrow (a + b) + c$$

$$a \times (b \times c) \longrightarrow (a \times b) \times c$$

$$a \times b + c \longrightarrow FMA(a, b, c)$$

Optimizations in Icing



Commutativity
(preserves IEEE754)

Example Optimizations:

$$a + b \longrightarrow b + a$$

$$a \times b \longrightarrow b \times a$$

$$a + (b + c) \longrightarrow (a + b) + c$$

$$a \times (b \times c) \longrightarrow (a \times b) \times c$$

$$a \times b + c \longrightarrow FMA(a, b, c)$$

Optimizations in Icing

Source

Target



s



t



Associativity
(no IEEE754)

Example Optimizations:

$$a + b \longrightarrow b + a$$

$$a \times b \longrightarrow b \times a$$

$$a + (b + c) \longrightarrow (a + b) + c$$

$$a \times (b \times c) \longrightarrow (a \times b) \times c$$

$$a \times b + c \longrightarrow FMA(a, b, c)$$

Optimizations in Icing

Source

Target



s



FMA introduction
(locally more accurate)

Example Optimizations:

$$a + b \longrightarrow b + a$$


$$a \times b \longrightarrow b \times a$$

$$a + (b + c) \longrightarrow (a + b) + c$$

$$a \times (b \times c) \longrightarrow (a \times b) \times c$$

$$a \times b + c \longrightarrow FMA(a, b, c)$$

Optimizations in Icing



Example Optimizations:

$$a + b \longrightarrow b + a$$

$$a \times b \longrightarrow b \times a$$

$$a + (b + c) \longrightarrow (a + b) + c$$

$$a \times (b \times c) \longrightarrow (a \times b) \times c$$

$$a \times b + c \longrightarrow FMA(a, b, c)$$

Floating-Point Values in Icing

IEEE754:

$$3.5 + 2.0 = 5.5$$

Icing:

Floating-Point Values in Icing

IEEE754:

Icing:

$3.5 + 2.0 = 5.5$ ← floating-point word

Floating-Point Values in Icing

IEEE754:

$3.5 + 2.0 = 5.5$ \longleftarrow floating-point word

Icing:

$3.5 + 2.0 =$


Floating-Point Values in Icing

IEEE754:

$3.5 + 2.0 = 5.5$ ← floating-point word

Icing:

$3.5 + 2.0 =$




Floating-Point Values in Icing

IEEE754:

$3.5 + 2.0 = 5.5$ ← floating-point word

Icing:

$3.5 + 2.0 =$



value tree for addition


Floating-Point Values in Icing

IEEE754:

$$3.5 + 2.0 = 5.5 \leftarrow \text{floating-point word}$$

Icing:

$$3.5 + 2.0 =$$



value tree for addition

$$3.5 + (2.0 + 1.5) = 12.25$$


Floating-Point Values in Icing

IEEE754:

$$3.5 + 2.0 = 5.5 \leftarrow \text{floating-point word}$$

Icing:

$$3.5 + 2.0 =$$



value tree for addition

$$3.5 + (2.0 + 1.5) = 12.25$$

$$3.5 + (2.0 + 1.5) =$$


Floating-Point Values in Icing

IEEE754:

$$3.5 + 2.0 = 5.5 \leftarrow \text{floating-point word}$$


Icing:

$$3.5 + 2.0 =$$



value tree for addition

$$3.5 + (2.0 + 1.5) = 12.25$$



Floating-Point Values in Icing


IEEE754:

$$3.5 + 2.0 = 5.5 \leftarrow \text{float}$$


Icing:



$$3.5 + (2.0 + 1.5) = 12.$$



value tree for addition



Icing's semantics

Allowed Optimization:

$$a \times b + c \longrightarrow FMA(a, b, c)$$

$$a \times b \longrightarrow b \times a$$

opt: (x * 2.4 + y)

Icing's semantics

Allowed Optimization:

$$a \times b + c \longrightarrow FMA(a, b, c)$$

$$a \times b \longrightarrow b \times a$$

} Included in the semantics

opt: (x * 2.4 + y)

Icing's semantics

Allowed Optimization:

$$a \times b + c \longrightarrow FMA(a, b, c)$$

$$a \times b \longrightarrow b \times a$$

} Included in the semantics

fine-grained control \longrightarrow opt: (x * 2.4 + y)

Icing's semantics


Allowed Optimization:

$a \times b + c \longrightarrow FMA(a, b, c)$

$a \times b \longrightarrow b \times a$

Included in the semantics

fine-grained control \longrightarrow opt: (x * 2.4 + y)



Icing's semantics


Allowed Optimization:

$a \times b + c \longrightarrow FMA(a, b, c)$

$a \times b \longrightarrow b \times a$

Included in the semantics

fine-grained control \longrightarrow opt: (x * 2.4 + y)



Icing's semantics

Allowed Optimization:


$a \times b + c \longrightarrow FMA(a, b, c)$

$a \times b \longrightarrow b \times a$

Included in the semantics

fine-grained control

Icing: a direct fit for fast-math
with fine-grained control and
support for different optimizations



Modelling the state-of-the-art

Unverified Compilers (gcc, clang,)

- aggressive optimizations
- no IEEE754 semantics
- no guarantees on the result

Verified Compilers (CakeML, ...)

- no floating-point optimizations
- IEEE754 semantics
- introduces no new behaviour

Icing provides:

greedy optimizer

IEEE754 Translator

Modelling the state-of-the-art

Unverified Compilers (gcc, clang,)

- aggressive optimizations
- no IEEE754 semantics
- no guarantees on the result

Verified Compilers (CakeML, ...)

- no floating-point optimizations
- IEEE754 semantics
- introduces no new behaviour

Icing provides:

greedy optimizer 

IEEE754 Translator 

What can we prove about the optimizers

Greedy optimizer:

⊢ if evaluating the greedily optimized program p returns v
then v is a possible result of evaluating p with the optimizations of the greedy optimizer

The greedy optimizer applies optimizations with respect to Icing semantics

IEEE754 translator:

⊢ after running the IEEE754 translator on program p no optimizations can be applied by Icing semantics

⊢ after running the IEEE754 translator on program p Icing semantics are deterministic no matter which optimizations are allowed

The IEEE754 translator preserves literal meaning (like CompCert/CakeML)

Distributivity in Icing

$$a \times (b + c) \longrightarrow a \times b + a \times c$$

Distributivity in Icing

$$a \times (b + c) \longrightarrow a \times b + a \times c$$

$$x * (y + z)$$

Distributivity in Icing

$$a \times (b + c) \longrightarrow a \times b + a \times c$$


$$x * (y + z)$$

Compiler

$$x * y + x * z$$

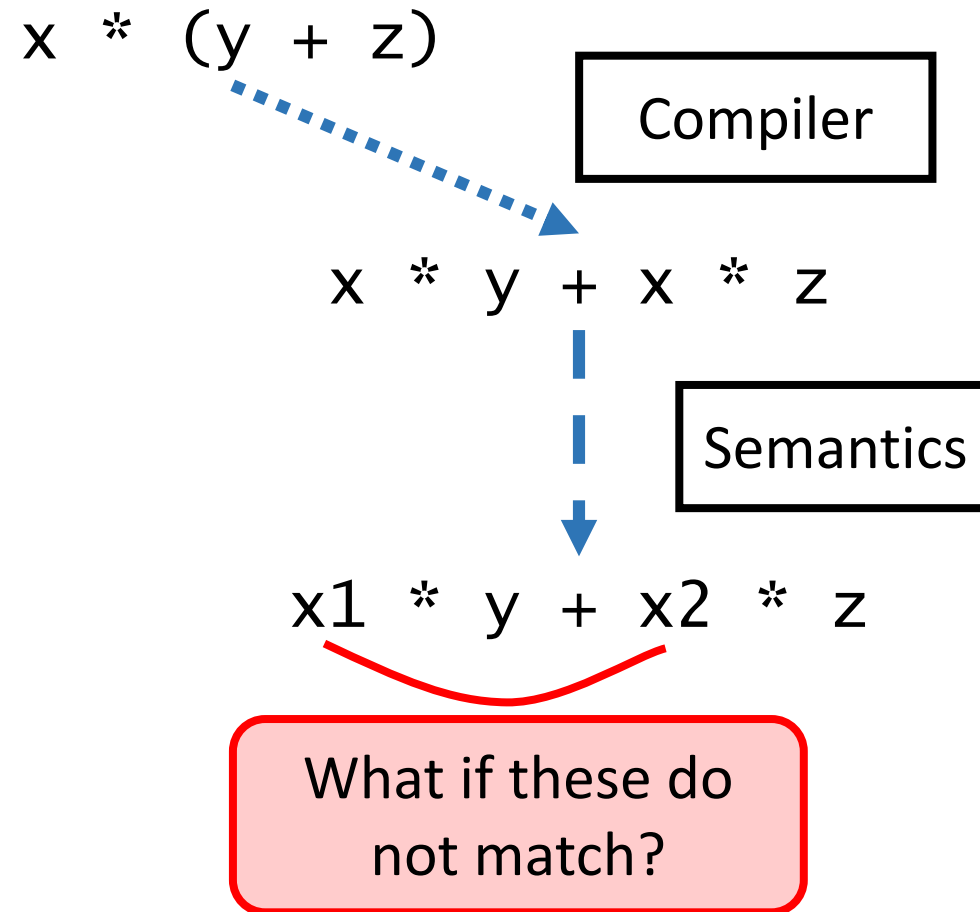
Distributivity in Icing

$$a \times (b + c) \longrightarrow a \times b + a \times c$$




Distributivity in Icing

$$a \times (b + c) \longrightarrow a \times b + a \times c$$



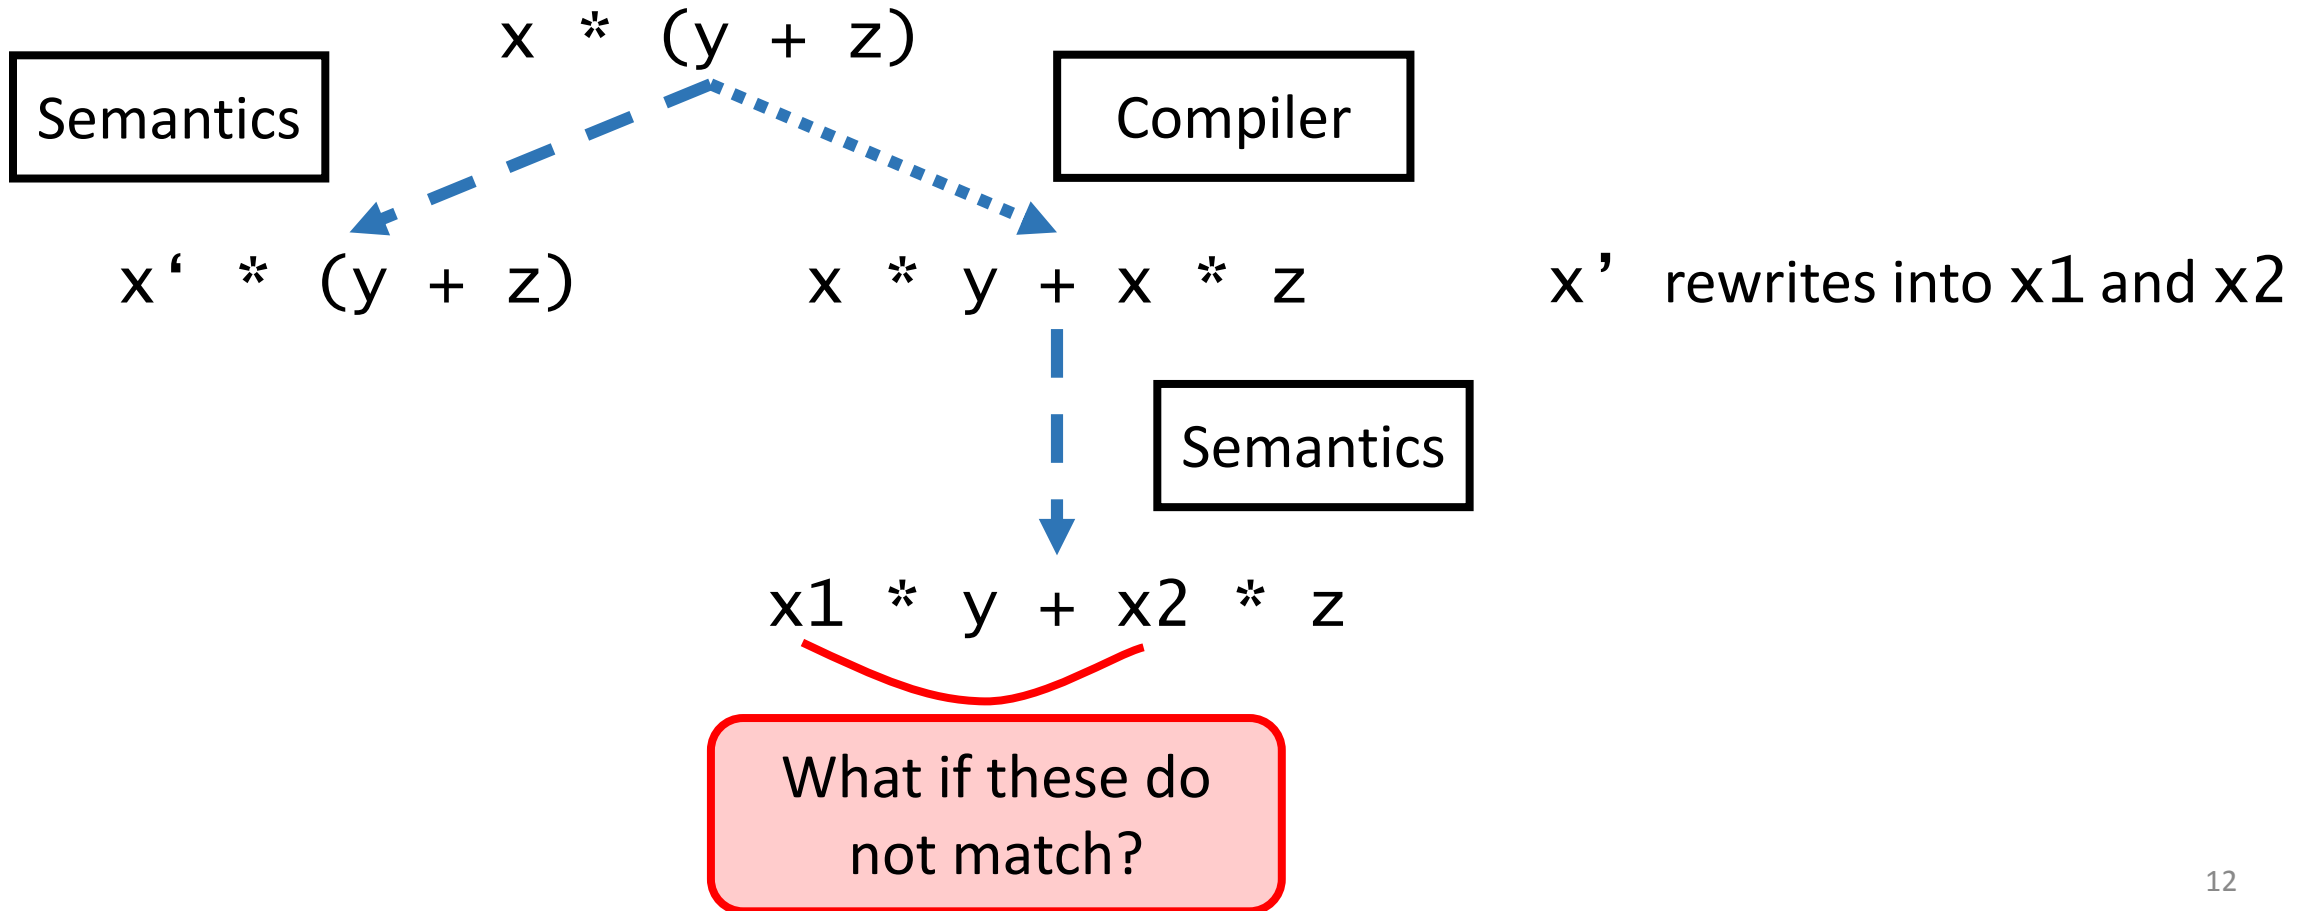
Distributivity in Icing

$$a \times (b + c) \longrightarrow a \times b + a \times c$$



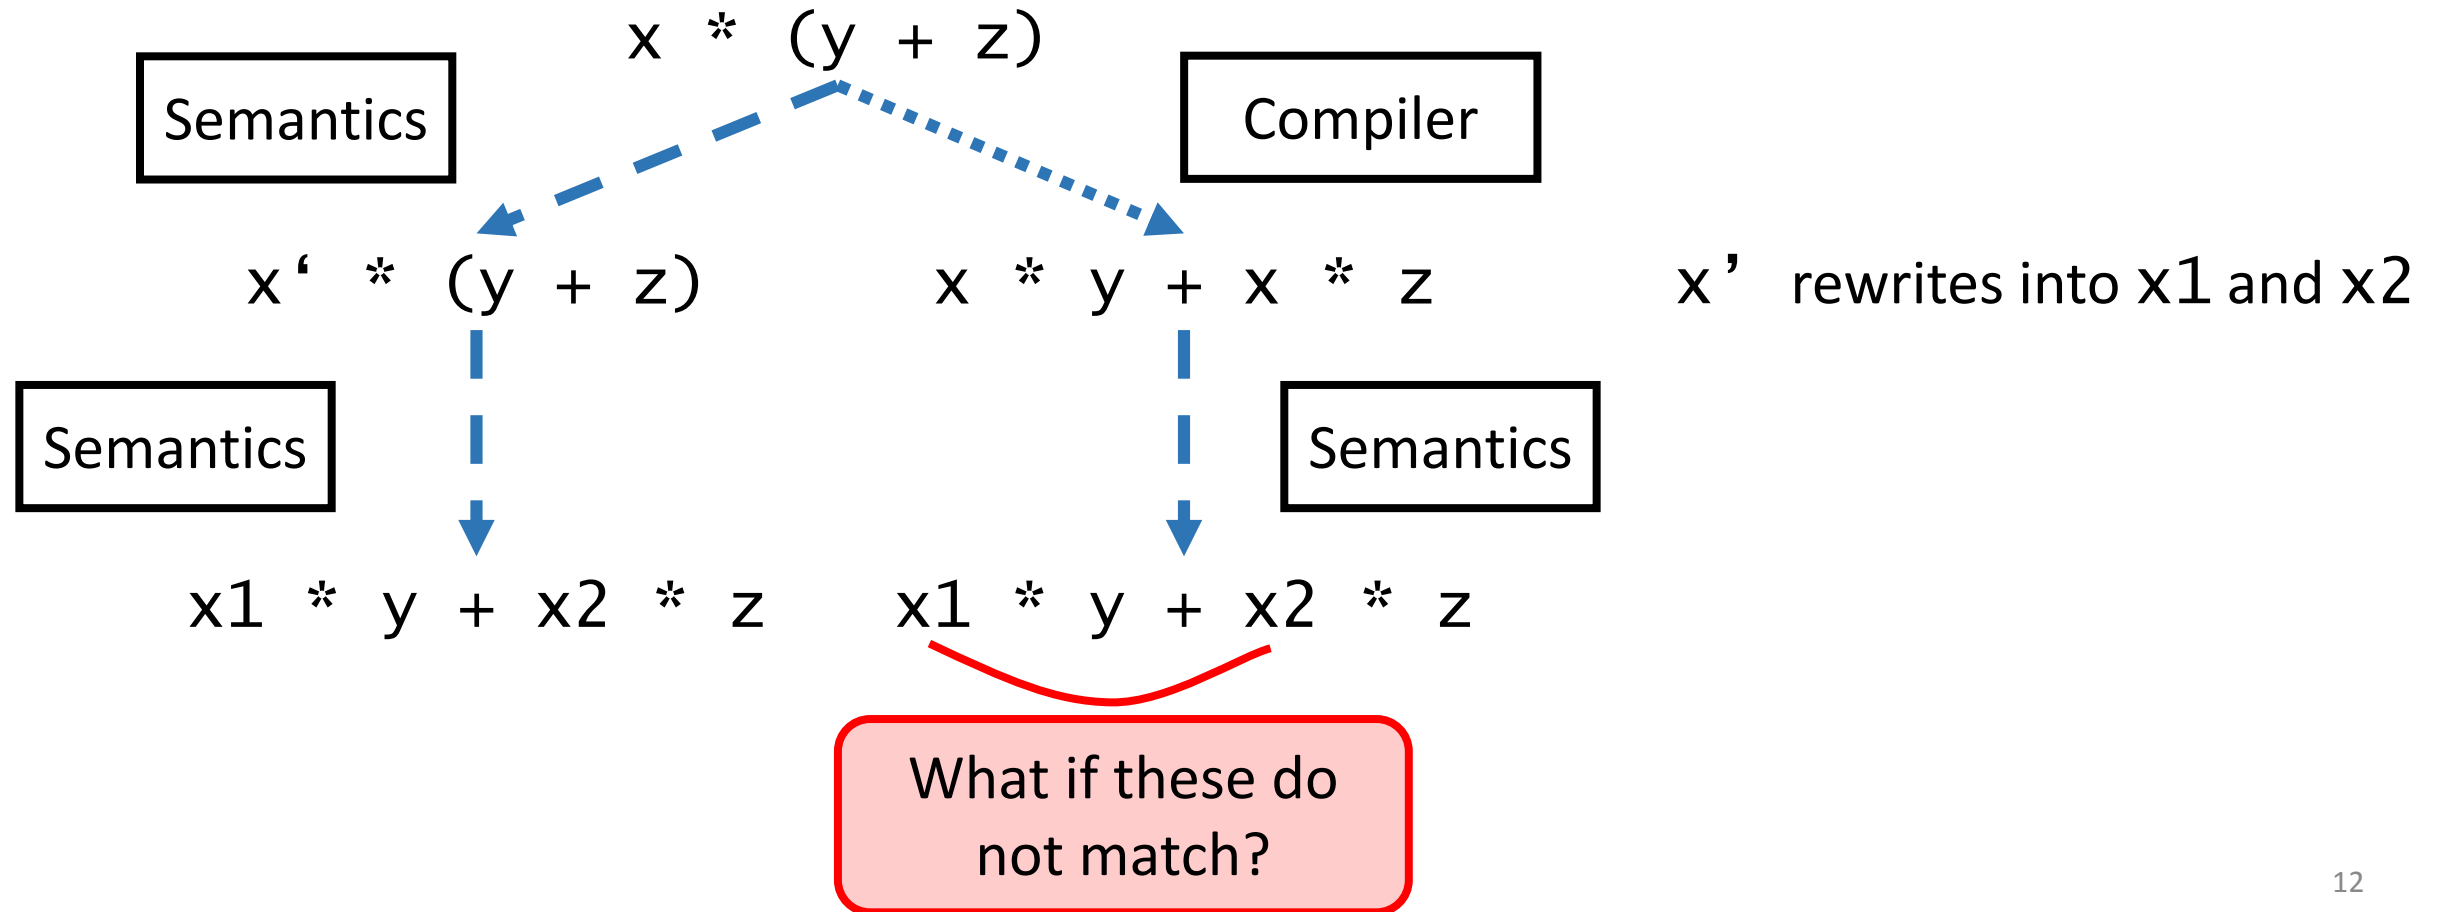
Distributivity in Icing

$$a \times (b + c) \longrightarrow a \times b + a \times c$$



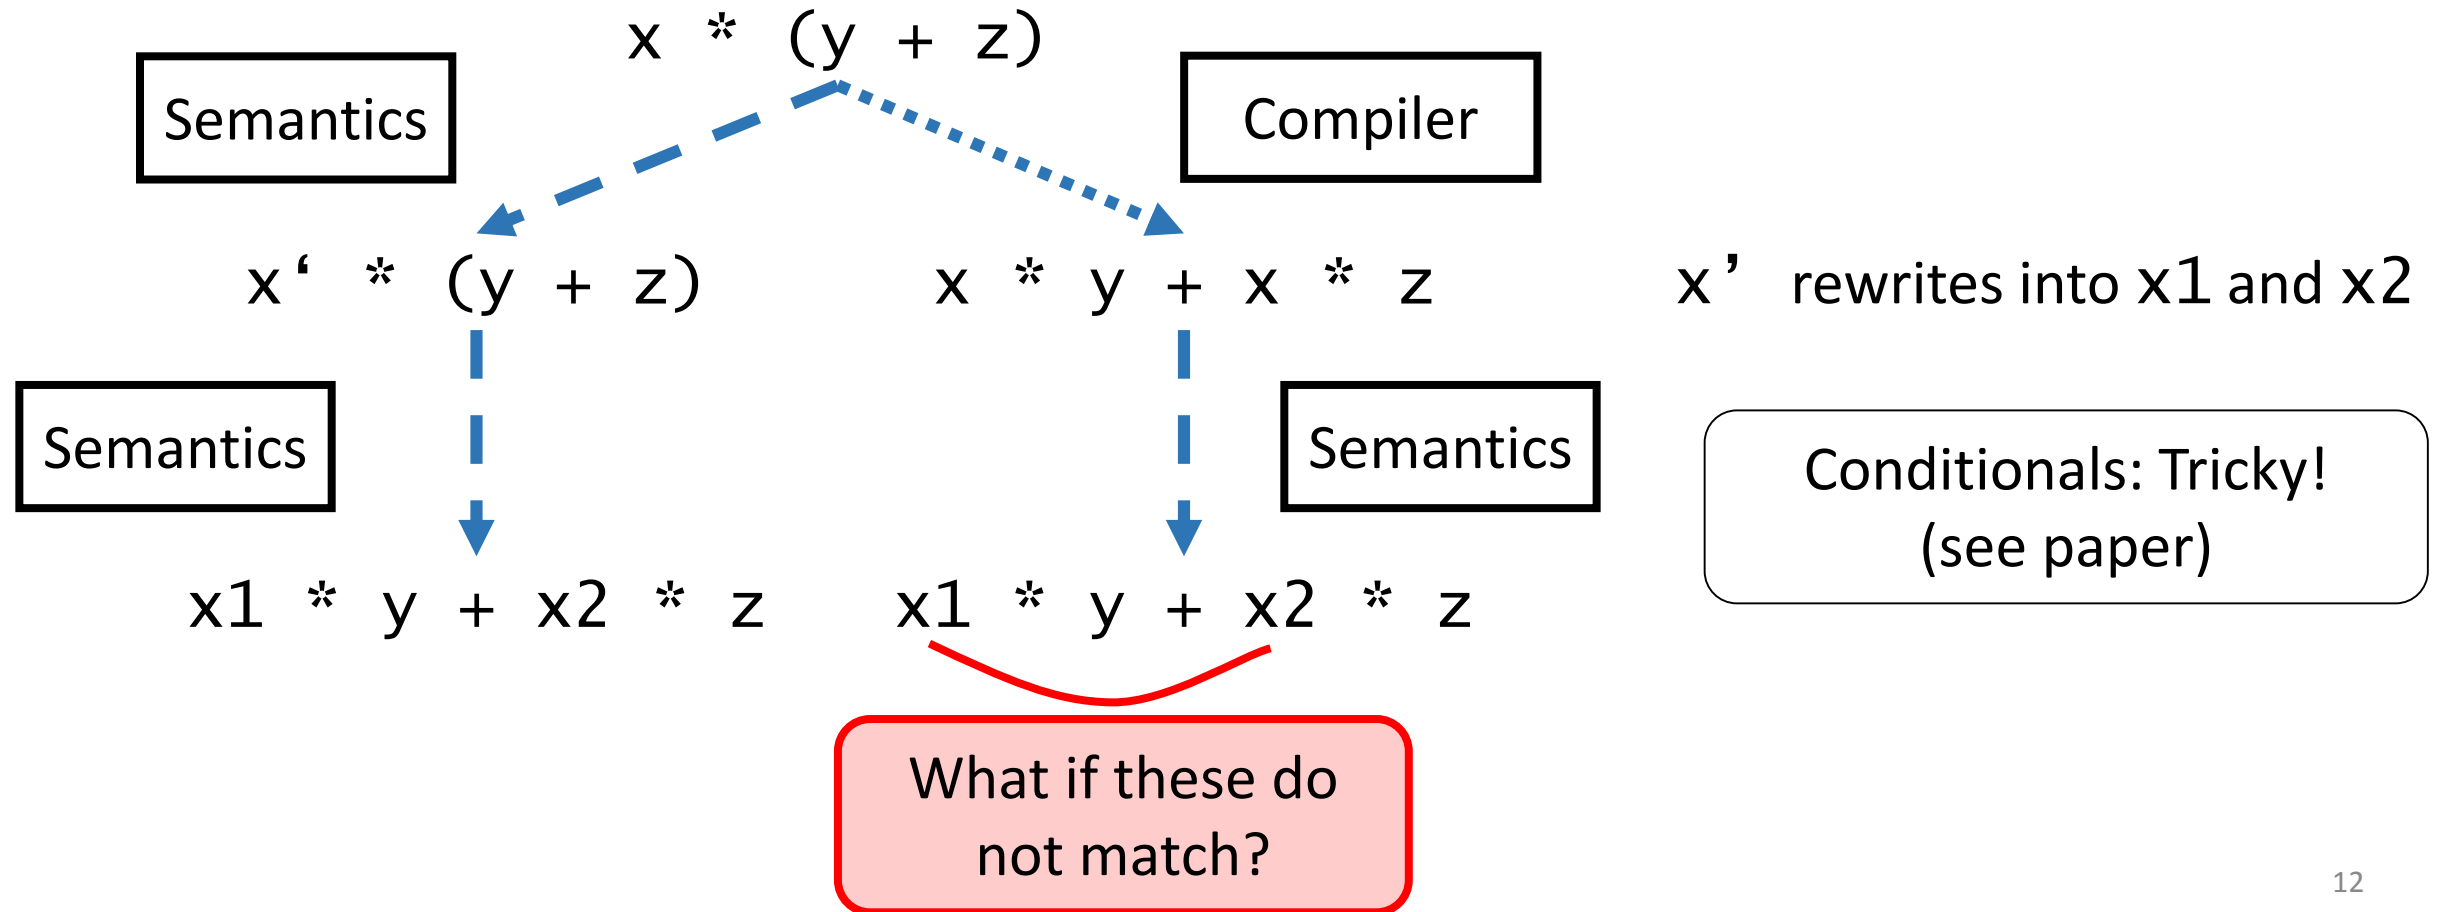
Distributivity in Icing

$$a \times (b + c) \longrightarrow a \times b + a \times c$$



Distributivity in Icing

$$a \times (b + c) \longrightarrow a \times b + a \times c$$



Handling more of gcc's rewrites

Enable fast-math mode. This defines the `__FAST_MATH__` preprocessor macro. Lossy assumptions about floating-point math. These include:

- Floating-point math obeys regular algebraic rules for real numbers (e.g. `* c == a * c + b * c`),
- operands to floating-point operations are not equal to `NaN` and `Inf`, and
- `+0` and `-0` are interchangeable.

Official clang documentation


Handling more of gcc's rewrites

Enable fast-math mode. This defines the `__FAST_MATH__` preprocessor macro. Lossy assumptions about floating-point math. These include:

- Floating-point math obeys regular algebraic rules for real numbers (e.g. `* c == a * c + b * c`),
- operands to floating-point operations are not equal to `NaN` and `Inf`, and
- `+0` and `-0` are interchangeable.

Official clang documentation

Handling more of gcc's rewrites




Enable fast-math mode. This defines the `__FAST_MATH__` preprocessor macro, which enables lossy assumptions about floating-point math. These include:

- Floating-point math obeys regular algebraic rules for real numbers (e.g. `* c == a * c + b * c`),
- operands to floating-point operations are not equal to `NaN` and `Inf`, and
- `+0` and `-0` are interchangeable.

Official clang documentation

Handling more of gcc's rewrites



gcc:


isNaN (c) \longrightarrow F

Enable fast-math mode. This defines the `FAST_MATH` preprocessor macro lossy assumptions about floating-point math. These include:

- Floating-point math obeys regular algebraic rules for real numbers (e.g. `* c == a * c + b * c`),
- operands to floating-point operations are not equal to NaN and Inf, and
- `+0` and `-0` are interchangeable.

Official clang documentation

Handling more of gcc's rewrites



Precondition allows to check condition
before applying a rewrite

Enable fast-math mode. This defines the `FAST_MATH` preprocessor macro lossy assumptions about floating-point math. These include:


- Floating-point math obeys regular algebraic rules for real numbers (e.g. `* c == a * c + b * c`),
- operands to floating-point operations are not equal to `NaN` and `Inf`, and
- `+0` and `-0` are interchangeable.

gcc:

`isNaN (c) → F`

Official clang documentation

Handling more of gcc's rewrites



Precondition allows to check condition before applying a rewrite


lcing:	gcc:
$c = c$	$\text{isNaN}(c) \rightarrow F$

Enable fast-math mode. This defines the `FAST_MATH` preprocessor macro lossy assumptions about floating-point math. These include:


- Floating-point math obeys regular algebraic rules for real numbers (e.g. `* c == a * c + b * c`),
- operands to floating-point operations are not equal to NaN and Inf, and
- `+0` and `-0` are interchangeable.

Official clang documentation

Handling more of gcc's rewrites



Precondition allows to check condition before applying a rewrite

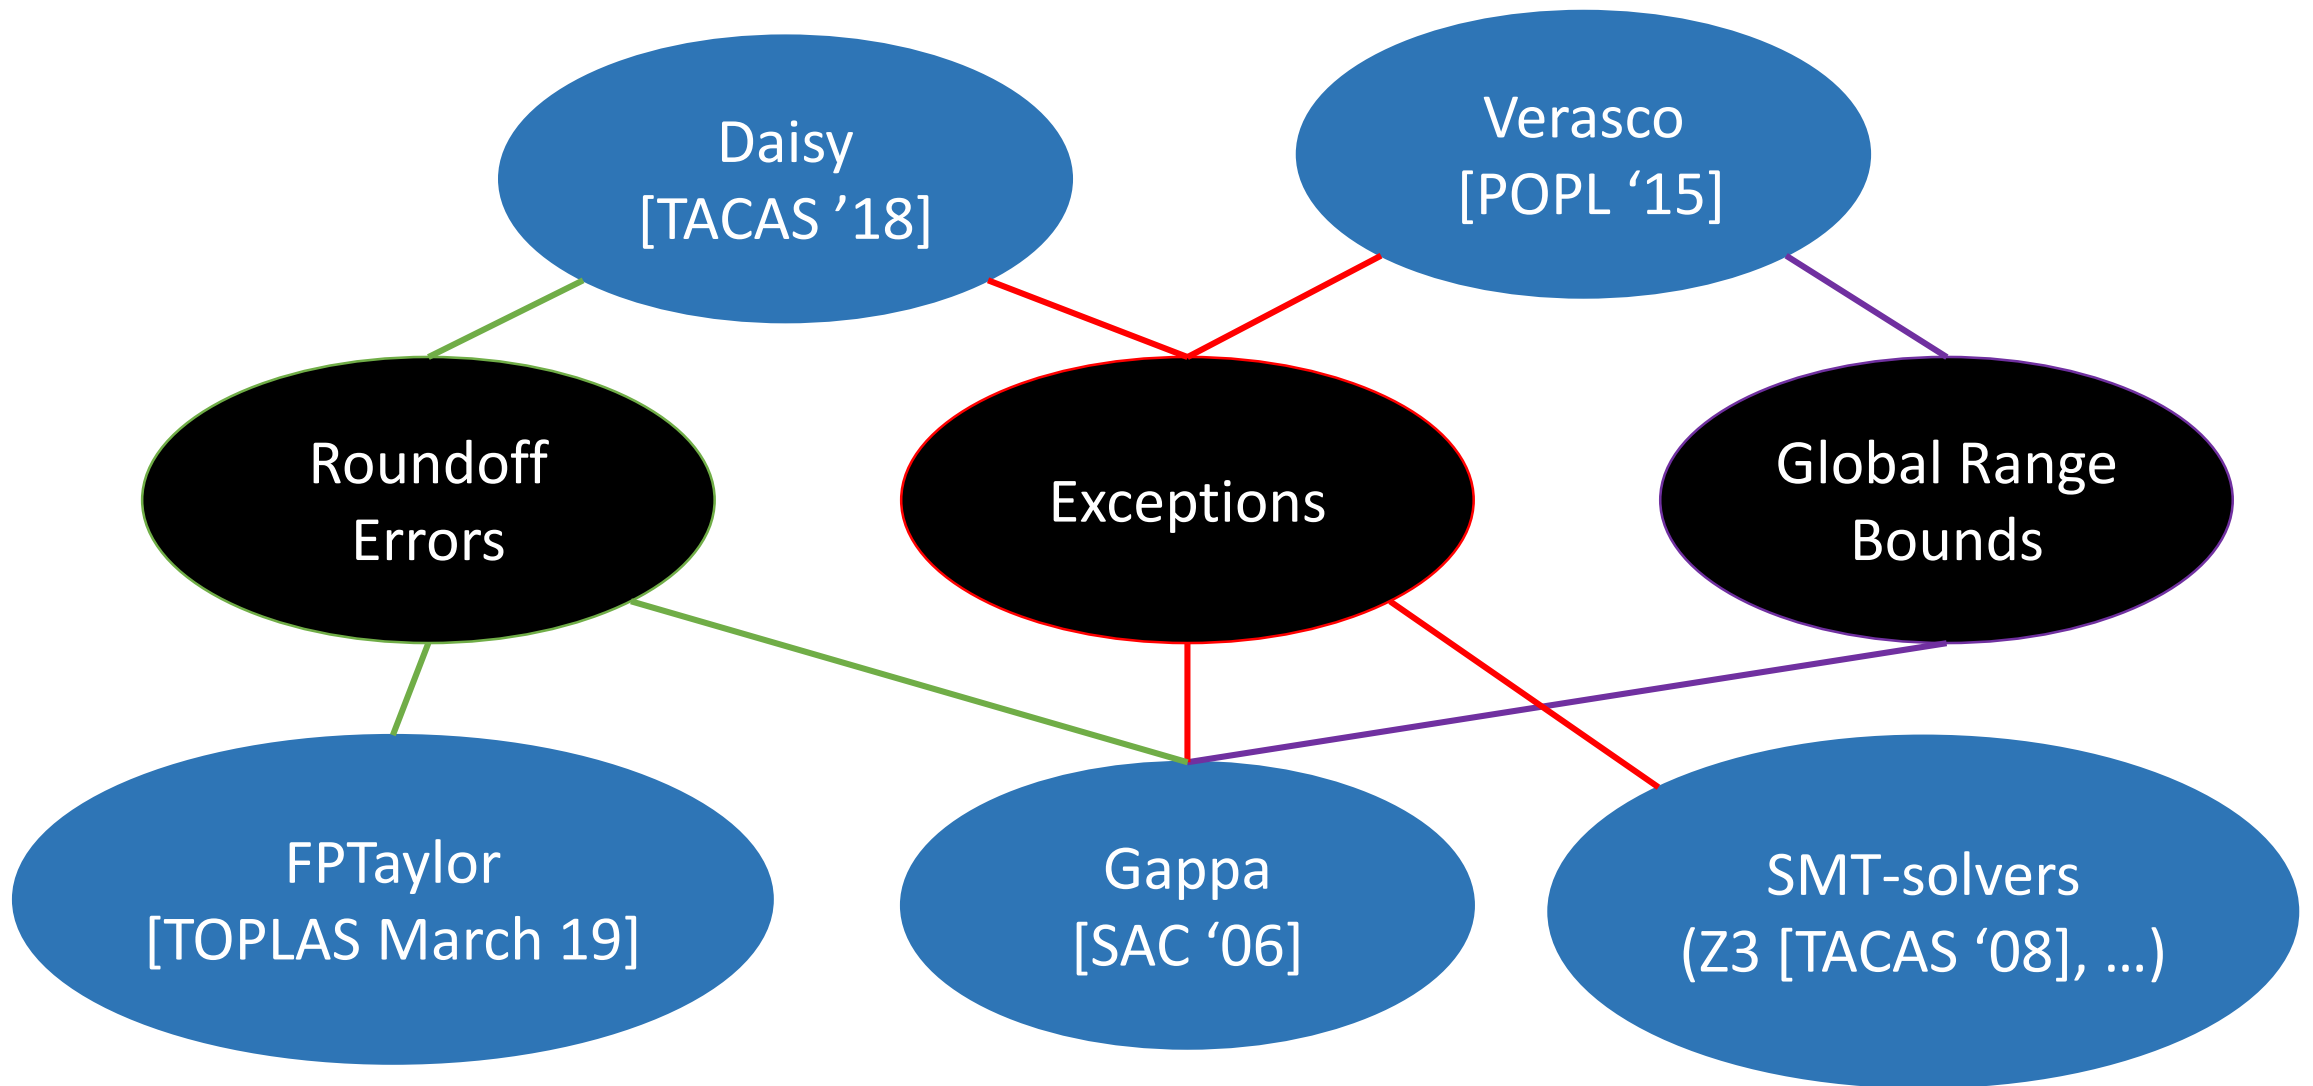


Enable fast-math mode. This defines the `FAST_MATH` preprocessor macro lossy assumptions about floating-point math. These include:

- Floating-point math obeys regular algebraic rules for real numbers (e.g. `* c == a * c + b * c`),
- operands to floating-point operations are not equal to NaN and Inf, and
- `+0` and `-0` are interchangeable.

Official clang documentation

How can the preconditions be checked



Icings interface to external tools




Discharge checks in-place

a, b, c
variables $\xrightarrow{\text{blue arrow}} a \times (b + c) \xrightarrow{\text{blue arrow}} a \times b + a \times c$

└ simple local check

\Rightarrow checked before applying optimization



Record assumed proposition

$c = c \xrightarrow{\text{blue arrow}} \text{isNaN}(c) \xrightarrow{\text{blue arrow}} \text{False}$

└ complex global property

\Rightarrow checked offline after compiling

What does gcc's fast-math actually do?

Nondeterministic Icing
(with optimizations)



deterministic Icing
(without optimizations)



CakeML source

Outlook:

- integrate with external tools
- verify larger optimizations
- integrate into CakeML semantics